



Build Data Pipelines with Lakeflow Spark Declarative Pipelines

Databricks Academy
May 2026



Agenda

Course Sections

- Introduction to Data Engineering in Databricks
- Lakeflow Spark Declarative Pipeline Fundamentals
- Building Lakeflow Spark Declarative Pipelines



Course Learning Objectives

- Understand the core concepts and components of Lakeflow Spark Declarative Pipelines (SDP), including the function and differences between streaming tables, materialized views, and temporary views.
- Identify and configure pipeline settings, such as compute, data assets, trigger modes, and advanced options.
- Develop a functional Spark Declarative Pipeline using the new pipeline editor and SQL-based syntax.
- Incorporate data quality expectations into a pipeline to validate and enforce data integrity.
- Analyze event logs and pipeline metrics to understand the full execution and lifecycle of a Spark Declarative Pipeline.
- Design and implement Change Data Capture (CDC) to a pipeline using AUTO CDC INTO to handle slowly changing dimensions (SCD).



Course Prerequisites (REQUIRED)



Fundamental Knowledge of the Databricks Platform

- Databricks Workspaces
- Apache Spark
- Unity Catalog Tables and the Medallion Architecture
- Unity Catalog



Knowledge of Ingesting Raw Data into a Table

- Familiarity with ingesting raw data source files using the **read_files** SQL function
 - CSV
 - JSON
 - TXT
 - Parquet



Experience Transforming Data with SQL

- Experience in writing **intermediate-level queries** using SQL
- Basic knowledge of **SQL Joins**



Lab Exercise Environment



Technical Details

- Your lab environment is provided by Vocareum.
- It will open in a new tab.
- It has been configured with the permissions and resources required to accomplish the tasks outlined in the lab exercise.
- Third party cookies must be enabled in your browser for Vocareum's user experience to work properly.
- Make sure to enable pop ups!





Introduction to Data Engineering in Databricks

Build Data Pipelines with Lakeflow Spark Declarative Pipelines

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Agenda

Section Overview – Introduction to Data Engineering in Databricks

- Data Engineering in Databricks
- What are Lakeflow Spark Declarative Pipelines?
- Course Setup and Creating a Pipeline
- Course Project Overview



Section Learning Objectives

- Understand the components of Data Engineering with Lakeflow in Databricks with Connect, Spark Declarative Pipelines, and Jobs.
- Explain how Lakeflow Spark Declarative Pipelines incrementally processes data using Streaming Tables and Materialized Views in batch or streaming Lakeflow Jobs.
- Demonstrate how to create a Lakeflow Spark Declarative Pipelines project in a Databricks Workspace.






Introduction to Data Engineering in Databricks

LECTURE

Data Engineering in Databricks



Data Engineering in Databricks




Data Engineering in Databricks

Simplifying batch and streaming ETL with automated reliability, optimizations and built-in data quality with **Lakeflow Spark Declarative Pipelines (SDP)**!



UNIFIED DATA ENGINEERING FOR THE DATA INTELLIGENCE PLATFORM

CONNECT



Efficient
ingestion
connectors

SPARK DECLARATIVE
PIPELINES



Accelerated ETL
development

JOBS



Reliable
orchestration for
analytics and AI

INDUSTRY LEADING DATA PROCESSING ENGINE
(Apache Spark + Structured Streaming)

UNIFIED GOVERNANCE



OPTIMIZED STORAGE





Introduction to Data Engineering in Databricks


LECTURE

What are Lakeflow Spark Declarative Pipelines?




What are Spark Declarative Pipelines?

It's difficult to build and operate reliable data pipelines




What are Spark Declarative Pipelines?

It's difficult to build and operate reliable data pipelines



What are Spark Declarative Pipelines?

Spark Declarative Pipelines: Reliable data pipelines made easy



What are Spark Declarative Pipelines?

Introducing Declarative Pipelines: Reliable data pipelines made easy

- **Simplified Pipeline Authoring**

Easily declare data ingestion and transformations with **SQL or Python**, and let Spark Declarative Pipelines handle the rest!

- **Intelligent Optimization at Scale**

Automated scaling and recovery improve reliability and reduce maintenance.

- **Unified Batch and Streaming**

Pipelines seamlessly adapt to **both near real-time and batch Lakeflow Jobs**, optimizing performance and cost.

The screenshot displays the Databricks interface for a declarative pipeline. On the left, a sidebar shows the 'Demo Pipeline' configuration, including a 'Last runs' section indicating a successful run in 1m 2s, and a 'Pipeline assets' section listing files like 'sample_exploration', 'sample_trips_demo_pipeline.sql', 'sample_zones_demo_pipeline.sql', and 'README.md'. The main area is divided into two panes. The left pane shows the SQL code for 'sample_trips_demo_pipeline.sql', which defines a materialized view by selecting pickup_zip and fare_amount from the nyctaxi.trips table. The right pane shows the 'Pipeline graph' with two nodes: 'sample_trips_demo_pipeline' (completed in 7s) and 'sample_zones_demo_pipeline' (completed in 4s). At the bottom, a 'Tables' section provides a summary of the pipeline's execution, including a table with columns for Name, Catalog, Schema, Type, Duration, Write, Expectations, Drop, Fail %, and Flows.

Name	Catalog	Schema	Type	Dura...	Writ...	Expectations	Drop...	Fail %	Flows
sample_trip...	labuser1043...	default	Materialized...	7s	-	Not defined	0	-	1
sample_zon...	labuser1043...	default	Materialized...	4s	-	Not defined	0	-	1



What are Spark Declarative Pipelines?

Connecting to Data Sources

Lakeflow Connect



File sources

(S3, ADLS, GCS, SFTP, Google Drive, Sharepoint, etc)



Message Queues

(Kafka, Kinesis, Pulsar, Pub/Sub, etc)



Databases

(SQL Server, PostgreSQL, MySQL, etc)



Software as a Service

(Workday, Salesforce, ServiceNow, Google Analytics, Oracle, Dynamics 365)

Spark Declarative Pipelines



Bronze



Silver




Gold



What are Spark Declarative Pipelines?


Simplifying Batch and Streaming ETL in the Medallion Architecture

Spark Declarative Pipelines




What are Spark Declarative Pipelines?

Overview of Incremental Processing in Declarative Pipelines




What are Spark Declarative Pipelines?

Overview of Incremental Processing in Declarative Pipelines (Run 1)





What are Spark Declarative Pipelines?

Overview of Incremental Processing in Declarative Pipelines (Run 2)



Creating a Spark Declarative Pipeline





Introduction to Data Engineering in Databricks

DEMONSTRATION

Course Setup and Creating a Pipeline



Notebook: 1 Demo – Course Setup and Creating a Pipeline



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](https://www.apache.org/).




Introduction to Data Engineering in Databricks

LECTURE

Course Project Overview



Course Project Overview





Lakeflow Spark Declarative Pipeline Fundamentals

**Build Data Pipelines with Lakeflow Spark Declarative
Pipelines**

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Agenda

Section Overview – Lakeflow Spark Declarative Pipeline Fundamentals

- Dataset Types Overview
- Simplified Pipeline Development
- Common Pipeline Settings
- Developing a Simple Pipeline
- Ensure Data Quality with Expectations



Section Learning Objectives

- Understand the core concepts and components of Lakeflow Spark Declarative Pipelines pipelines, including how streaming tables, materialized views, and temporary views function and differ.
- Identify and configure pipeline settings such as compute, data assets, trigger modes, and advanced options.
- Develop a functional Declarative Pipeline using the new pipeline editor and SQL-based syntax for the pipeline.
- Incorporate data quality expectations into a Declarative Pipeline pipeline to validate and enforce data integrity.






Lakeflow Spark Declarative Pipeline Fundamentals

LECTURE

Dataset Types Overview



Dataset Types Overview

How streaming tables, materialized views, and views process data



STREAMING TABLE (ST)

A table with support for streaming or incremental data processing, **only processing new data**



MATERIALIZED VIEW (MV)



VIEWS



Dataset Types Overview

STREAMING TABLE Overview

STREAMING TABLE

- Support for **batch** or **streaming** for **incremental** data processing (exactly **once processing** of the files)
- Each time an streaming table is refreshed, data added to the source tables is **appended** to the **streaming table**
- SQL code that creates streaming tables:
CREATE OR REFRESH STREAMING TABLE
- In SQL, you must use **FROM STREAM read_files()** to **invoke Auto Loader functionality** for incremental streaming reads and checkpoints
- File names are **guaranteed to be read only once**



Dataset Types Overview

Create a STREAMING TABLE with SQL

Create a STREAMING TABLE named **orders_bronze** from **JSON** files

```
CREATE OR REFRESH STREAMING TABLE 1_bronze_db.orders_bronze AS
```

```
SELECT
```

```
*,
```

```
current_timestamp() AS processing_time,
```

```
_metadata.file_name AS source_file
```

```
FROM STREAM read_files(
```

```
  "/Volumes/dbacademy/ops/labuser/orders",
```

```
  format => 'JSON');
```

Creates a streaming table named **orders_bronze** in the **1_bronze_db** schema

Use the **SELECT** clause to specify the columns to read into the streaming table from your data source

- The **STREAM** keyword indicates to use streaming read for the source files (JSON)
- The **read_files()** function reads files under a provided location and returns the data in tabular form

Note: The query is using the default catalog



Dataset Types Overview

Create a STREAMING TABLE with SQL

Create a STREAMING TABLE named **orders_silver** that reads from the **STREAMING orders_bronze** table

```
CREATE OR REFRESH STREAMING TABLE 2_silver_db.orders_silver AS
```

```
SELECT
```

```
    order_id,  
    timestamp(order_timestamp) AS order_timestamp,  
    customer_id,  
    notifications
```

Specify your SQL transformations for the streaming data from the bronze table

```
FROM STREAM 1_bronze_db.orders_bronze;
```

Use STREAM keyword with the name of the STREAMING TABLE you are reading from to transform new rows

Note: The query is using the default catalog




Dataset Types Overview

How streaming tables, materialized views, and views process data



STREAMING TABLE (ST)

A table with support for streaming or incremental data processing, **only processing new data**



MATERIALIZED VIEW (MV)

- Records are processed as required to return accurate results for the **current data state**
- Used for data processing tasks such as:
 - Transformations
 - Aggregations
 - Pre-computing slow queries
 - Frequently used computations



VIEWS




Dataset Types Overview

MATERIALIZED VIEW Overview

MATERIALIZED VIEW

- Each time a materialized view is updated, **query results are recalculated to reflect changes in upstream datasets**
- Created and **kept up-to-date** by pipeline
- Use the **CREATE OR REFRESH MATERIALIZED VIEW** syntax
- Can be used **anywhere in your pipeline**, not just in the gold layer
- **Where applicable**, results are **incrementally refreshed for materialized views**, avoiding the need to completely rebuild the materialized view when new data arrives (**Serverless Only**)
- **Incremental refresh for materialized views** is a cost-based optimizer to power fast and efficient transformations for materialized views on Serverless compute



Dataset Types Overview

Create a MATERIALIZED VIEW with SQL

Create a MATERIALIZED VIEW named **gold_orders_by_date** that summarizes the **orders_silver** table

```
CREATE OR REFRESH MATERIALIZED VIEW 3_gold_db.gold_orders_by_date AS
```

```
SELECT
```

```
  date(order_timestamp) AS order_date,
```

```
  count(*) AS total_daily_orders
```

```
FROM 2_silver_db.orders_silver
```

```
GROUP BY date(order_timestamp);
```

Create the materialized view **gold_orders_by_date** in the **3_gold_db** schema (database)

Specify the the source table without the STREAM keyword to utilize the source STREAMING table

NOTE: Where possible, queries will be incrementally refreshed

Note: The query is using the default catalog




Dataset Types Overview

How streaming tables, materialized views, and views process data



STREAMING TABLE (ST)

A table with support for streaming or incremental data processing, **only processing new data**



MATERIALIZED VIEW (MV)

- Records are processed as required to return accurate results for the **current data state**
- Used for data processing tasks such:
 - Transformations
 - Aggregations
 - pre-computing slow queries
 - frequently used computations.



VIEWS

Constructs a virtual table with no physical data based on the query in your Declarative Pipelines

- TEMPORARY VIEW
- VIEW



Dataset Types Overview

TEMPORARY VIEW Overview

TEMPORARY VIEW

- Temporary views are only **persisted across the lifetime** of the pipeline and **private to the defining pipeline**
- They are **not registered** as an object to Unity Catalog
- Use the **CREATE TEMPORARY VIEW** statement
- TEMPORARY VIEWS are useful as **intermediate queries** that **are not exposed** to end users



Dataset Types Overview


VIEW Overview

VIEW

- Constructs a virtual table with no physical data based on the result-set of a SQL query in your pipelines
- They are **registered** as an object to Unity Catalog
- Use the **CREATE VIEW** statement


LIMITATIONS

- The pipeline must be a Unity Catalog pipeline
- Views cannot have streaming queries, or be used as a streaming source for a pipeline



Dataset Types Overview

Existing users of DLT will notice that we've evolved the names



The **semantics remain the same**, and **we'll support the old syntax** for compatibility.
Our goal is to simplify the syntax and match other systems



Dataset Types Overview

Pipeline dependencies are parsed automatically, code order is not required

Raw JSON → Bronze


```
CREATE OR REFRESH STREAMING TABLE
1_bronze_db.orders_bronze AS
SELECT
*,
current_timestamp() AS processing_time,
_metadata.file_name AS source_file
FROM STREAM read_files(
  "/Volumes/dbacademy/ops/labuser/orders",
  format => 'JSON');
```

Bronze → Silver

```
CREATE OR REFRESH STREAMING TABLE 2_silver_db.orders_silver
AS
SELECT
  order_id,
  timestamp(order_timestamp) AS order_timestamp,
  customer_id,
  notifications
FROM STREAM 1_bronze_db.orders_bronze;
```

Silver → Materialized View

```
CREATE OR REFRESH MATERIALIZED VIEW 3_gold_db.gold_orders_by_date
AS
SELECT
  date(order_timestamp) AS order_date,
  count(*) AS total_daily_orders
FROM 2_silver_db.orders_silver
GROUP BY date(order_timestamp);
```






Lakeflow Spark Declarative Pipeline Fundamentals

LECTURE

Simplified Pipeline Development



Simplified Pipeline Development

Multi-file editor in Lakeflow Spark Declarative Pipelines

What's New

1. Pipeline asset browser
2. Multi-file code editor with features for step-by-step development
3. Pipeline-specific toolbars
4. Interactive DAG
5. Data previews
6. Execution insights panels
 - Easier debugging
 - Faster validation with dry run

With Spark Declarative Pipelines, you can write code using .sql or .py files (recommended), or use notebooks.

As of Data & AI Summit 2025






Lakeflow Spark Declarative Pipeline Fundamentals

LECTURE

Common Pipeline Settings



Common Pipeline Settings

Overview

The screenshot displays the Databricks Lakeflow Spark Declarative Pipelines editor. On the left, a sidebar shows the pipeline configuration and a list of files. A red box highlights the 'Pipeline configuration' tab, and a red arrow points to the 'Settings' button. A large red text box in the center states: "You can edit Lakeflow Spark Declarative Pipelines settings directly in the editor". The main editor area shows the SQL code for the pipeline, which includes a streaming table definition and a query. Below the code, a table lists the pipeline's components:


Name	Catalog	Schema	Type	Duration	Written rec...	Expected
orders_bronze	dbacademy_peter_styliadis	1_bronze_db	Streaming table	7s	174	
orders_silver	dbacademy_peter_styliadis	2_silver_db	Streaming table	3s	174	
gold_orders_by_date	dbacademy_peter_styliadis	3_gold_db	Materialized view	5s	7	

On the right, a 'Pipeline settings' panel is shown, detailing the pipeline's configuration:

- Pipeline settings**
 - Pipeline ID: ccb9eaa8-dfe2-4925-a36c-75e3c98b8a30
 - Pipeline type: ETL pipeline
 - Pipeline name: 1 - Developing a Simple DLT Pipeline Project
 - Creator: peter.styliadis@databricks.com
 - Run as: peter.styliadis@databricks.com
- Code assets**
 - Root folder: /Workspace/Users/p... DLT Pipeline Project
 - Source code: /Workspace/Users/p...s/orders_pipeline.sql
 - Configure paths
- Default location for data assets**
 - Default catalog: dbacademy_peter_styliadis
 - Default schema: default
 - Edit catalog and schema
- Compute**
 - Serverless
- Configuration**
 - source: /Volumes/dbacademy/ops/dbacademy_peter_styliadis
 - Edit configuration
- Tags**
 - Add tags
- Budget**



Common Pipeline Settings



Compute




Code Assets



Configuration
(parameters)



Common Pipeline Settings



Compute



Code Assets



Configuration
(parameters)


SERVERLESS

- Databricks **recommends** developing new pipelines using serverless
- **Optimizes** cost while **maintaining** performance
- **Enables** you to focus on code rather than managing infrastructure
- Optional **Serverless Performance optimized** setting for time-sensitive Lakeflow Jobs

- Serverless compute includes **Incremental refresh for materialized views**
- A cost-based optimizer to power fast and efficient transformations for materialized views on Serverless compute



Common Pipeline Settings



Compute

SERVERLESS

- Databricks **recommends** developing new pipelines using serverless
- **Optimizes** cost while **maintaining** performance
- **Enables** you to focus on code rather than managing infrastructure
- Optional **Serverless Performance optimized** setting for time-sensitive Lakeflow Jobs



Code Assets

CLASSIC (fixed size)

- Users need **permission** to create compute for Declarative Pipelines
- Workspace admins **can configure cluster policies** to provide users with access to compute resources for Declarative Pipelines



Configuration (parameters)

AUTOSCALING

- Enhanced autoscaling is **enabled by default** for all new pipelines.
- **Optimizes cluster utilization** by automatically allocating cluster resources based on workload volume



Common Pipeline Settings



Compute



Code Assets



Configuration
(parameters)

Pipeline settings UI JSON ×

Pipeline ID5f9c6978-8d8b-4c21-b644-f861fb050656

Pipeline typeETL pipeline

Pipeline nameNew Pipeline 2025-04-10 10:31

Creatorpeter.styliadis@databricks.com

Run aspeter.styliadis@databricks.com

Code assets

Root folder ⓘ

/Workspace/Users/p...ith DLT Expectations

Source code ⓘ

/Workspace/Users/p...s/orders_pipeline.sql

Configure paths

Default location for data assets

Default catalog ⓘdbacademy_peter_styliadis

Default schema ⓘdefault

Edit catalog and schema

← Pipeline source code

Pipeline root folder

Default location for source code. Automatically appended to Python sys.path to support imports of modules and packages in this folder from source code files.

/Workspace/Users/peter

com/build-data-pipeline

Source code

SQL and Python files within specified folders or provided as individual file paths are processed when the pipeline runs.

/Workspace/Users/peter

com/build-data-p


+ Add path

The **Pipeline root folder** automatically includes all relevant files within the folder for the pipeline project (can be a Git folder)

The **source code** section references include subfolders (or files) for the pipeline (**.py, .sql, notebooks**)



Common Pipeline Settings



Compute



Code Assets



Configuration
(parameters)

A pipeline's configuration is **a map of key value pairs** that can be used to parameterize your code:

- Improve code readability/maintainability
- Reuse common parameters in multiple pipelines files

← Configuration

Key*	Value	
source	/Volumes/dbacademy/ops/dbac	🗑️
<button>Add parameter</button>		

```
CREATE STREAMING TABLE bronze
SELECT *
FROM STREAM read_files(
  "${source}/orders",
  format => 'JSON');
```

In SQL reference the key's value using the **`${key-name}`** syntax



Common Pipeline Settings

Common Settings

There are **various other settings** we will explore throughout the course, including:

- Pipeline settings
- Code assets
- Default location for data assets
- Compute
- Configuration
- Environment
- Tags
- Budget
- Advanced settings

The screenshot shows the 'Pipeline settings' dialog in Databricks. The 'UI' tab is selected. The settings are organized into several sections, each highlighted with a red box in the original image:

- Pipeline settings:** Contains fields for Pipeline ID (ccb9eaa8-dfe2-4925-a36c-75e3c98b8a30), Pipeline type (ETL pipeline), Pipeline name (1 - Developing a Simple DLT Pipeline Project), Creator (peter.styliadis@databricks.com), and Run as (peter.styliadis@databricks.com).
- Code assets:** Contains fields for Root folder (/Workspace/Users/p... DLT Pipeline Project) and Source code (/Workspace/Users/p...s/orders_pipeline.sql), along with a 'Configure paths' button.
- Default location for data assets:** Contains fields for Default catalog (dbacademy_peter_styliadis) and Default schema (default), along with an 'Edit catalog and schema' button.
- Compute:** Contains a 'Serverless' button.
- Configuration:** Contains a 'source' field with the value /Volumes/dbacademy/ops/dbacademy_peter_styliadis, along with an 'Edit configuration' button.
- Tags:** Contains an 'Add tags' button.
- Budget:** This section is partially visible at the bottom.






Lakeflow Spark Declarative Pipeline Fundamentals

DEMONSTRATION

Developing a Simple Pipeline



Notebook: 2 Demo – Developing a Simple Pipeline




© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](https://www.apache.org/).



Lakeflow Spark Declarative Pipeline Fundamentals

LECTURE

Ensure Data Quality with Expectations



Ensure Data Quality with Expectations

Creating Expectations

With Lakeflow Spark Declarative Pipelines you can apply **data quality rules** to validate rows during ETL processes to gain insights into data integrity using the following

```
CONSTRAINT constraint-name EXPECT (column condition) [ON VIOLATION action]
```

Action	SQL Syntax Example	Result
WARN (default)	CONSTRAINT valid_notification EXPECT (notifications IN ('Y','N'))	- Invalid rows are still written to the target - Logs include counts of valid vs. invalid records and other metrics
DROP	CONSTRAINT valid_date EXPECT (order_timestamp > "2021-01-01") ON VIOLATION DROP ROW	- Invalid rows dropped from table - The count is logged alongside other metrics
FAIL	CONSTRAINT valid_id EXPECT (customer_id IS NOT NULL) ON VIOLATION FAIL UPDATE	- Causes a failure of a single flow and does not cause other flows in your pipeline to fail. Manual intervention is required

Ensure Data Quality with Expectations

Adding Data Quality Expectations Example

```
CREATE OR REFRESH STREAMING TABLE 2_silver_db.orders_silver
```

```
(
```

```
  CONSTRAINT valid_notifications EXPECT (notifications IN ('Y', 'N')),
```

```
  CONSTRAINT valid_date EXPECT (order_timestamp > "2021-01-01") ON VIOLATION FAIL UPDATE,
```

```
  CONSTRAINT valid_id EXPECT (customer_id IS NOT NULL) ON VIOLATION DROP ROW
```

```
)
```

WARN in metrics

DROP the row

FAIL THE PIPELINE

```
AS
```

```
SELECT
```

```
  Order_id,
```

```
  timestamp(order_timestamp) AS order_timestamp,
```

```
  Customer_id,
```


```
  notifications
```

```
FROM STREAM 1_bronze_db.orders_bronze;
```



Ensure Data Quality with Expectations

Actions Overview



As of 2025Q2 materialized views that use **expectations** are **always fully refreshed**






Lakeflow Spark Declarative Pipeline Fundamentals

DEMONSTRATION

Adding Data Quality Expectations



Notebook: 3 Demo – Adding Data Quality Expectations




© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Lakeflow Spark Declarative Pipeline Fundamentals

LAB EXERCISE

Create a Pipeline



Notebook: 4 Lab – Create a Pipeline



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Building Lakeflow Spark Declarative Pipelines

**Build Data Pipelines with Lakeflow Spark Declarative
Pipelines**

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Agenda

Section Overview – Building Pipelines

- **Streaming Joins Overview**
- **Deploying a Pipeline to Production**
- **Change Data Capture (CDC) Overview**
- **Change Data Capture with AUTO CDC INTO**
- **Additional Features Overview**



Section Learning Objectives

- Deploy a Lakeflow Spark Declarative Pipeline in production by modifying configuration options like mode, schedule, email notifications and more.
- Analyze event logs and pipeline metrics to examine the entirety of a pipeline.
- Design and implement a Change Data Capture (CDC) Declarative Pipeline using AUTO CDC INTO to handle slowly changing dimensions (SCD).





Building Lakeflow Spark Declarative Pipelines

LECTURE

Streaming Joins Overview



Streaming Joins Overview

Stream-Snapshot Join Overview

Streaming Table to Static Table




The goal is to **incrementally** join new data with a **static table**



Streaming Joins Overview

Stream-Snapshot Join Overview

Streaming Table to Static Table




The new data in the streaming table JOINS using all of the data in the lookup table



Streaming Joins Overview

Stream-Snapshot Join Overview

Streaming Table to Static Table



Only the new data in the streaming table JOINS using all of the data in the lookup table




Streaming Joins Overview

Joining Streaming Tables with a Materialized View

Streaming Table to Static Table

Streaming Table to Streaming Table with a Materialized View



The goal is to take **all rows from two streaming tables** and join them together each time the pipeline is run




Streaming Joins Overview

Joining Streaming Tables with a Materialized View

Streaming Table to Static Table

Streaming Table to Streaming Table with a Materialized View



The materialized view **efficiently computes all the data** in the streaming tables and joins the rows.




Streaming Joins Overview

Joining Streaming Tables with a Materialized View

Streaming Table to Static Table

Streaming Table to Streaming Table with a Materialized View



As new data is added to the streaming tables, the materialized view will again efficiently compute all the data in the streaming tables and joins the rows



Streaming Joins Overview


Stream-Stream Joins

Streaming Table to Static Table

Streaming Table to Streaming Table with a Materialized View

Outside the scope of the course

Incremental Stream to Stream Joins



The goal is to **incrementally** join new data from two tables, **past data is not used**

[Databricks Streaming and DLT](#) (Advanced)
Course






Building Lakeflow Spark Declarative Pipelines

LECTURE

Deploying a Pipeline to Production



Deploying a Pipeline to Production



Deploying a Pipeline to Production

Common Production Tasks



Scheduling



Email
Notifications



Monitor
pipelines



Deploying a Pipeline to Production

Common Production Tasks



Scheduling



Email
Notifications



Monitor
pipelines

Triggered

- Can be triggered **manually** or run on a **schedule**
- **Refreshes** selected tables using data available at the start, then stops once the update is complete

Continuous

- **Continuously** processes new data to keep streaming tables and materialized views up to date in **near real time**
- Monitors dependencies and **updates only** when source data changes



Deploying a Pipeline to Production

Common Production Tasks



Scheduling



Email
Notifications



Monitor
pipelines

You can configure **email notifications** when **scheduling** the pipeline

Schedule

Every at :

☐ Show cron syntax

Timezone

Performance optimization

☐ Performance optimized New

More options

Notifications

☒ Start ☒ Success ☒ Failure



Deploying a Pipeline to Production

Common Production Tasks



Scheduling




Email
Notifications



Monitor
pipelines

- The **pipeline event log** captures all key information about the pipeline:

- Audit logs
- Data Quality Checks
- Pipeline Progress
- Data Lineage



Deploying a Pipeline to Production

Querying the Declarative Pipeline Event Log

- **Publish the event log** as a **Delta table** using the **advanced settings**
 - Specify the table **location** (catalog, schema)
 - Define the **table** name

Query the pipeline event log table

```
SELECT * FROM <catalog>.<schema>.<event_log_table_name>
```



Deploying a Pipeline to Production

Querying the Declarative Pipeline Event Log

- Publish the event log as a Delta table using the advanced settings
 - Specify the table location (catalog, schema)
 - Define the table name
- (DEFAULT) The Event Log is written as a **hidden Delta table**
 - Located in the pipeline **default catalog** and **schema**
 - View the [query the event log documentation](#) to access the hidden event log





Building Lakeflow Spark Declarative Pipelines

DEMONSTRATION

Deploying Your Pipeline to Production



Notebook: 5 Demo – Deploying a Pipeline to Production



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Building Lakeflow Spark Declarative Pipelines

LECTURE


Change Data Capture (CDC) Overview



Change Data Capture (CDC) Overview

What is Change Data Capture?

- Change Data Capture (CDC) is a technique **used to track and capture changes in a data source** (such as a database, lakehouse or data warehouse).
- Then **applying those changes** into the Lakehouse or data warehouse



Change Data Capture (CDC) Overview

Slowing Changing Dimensions (SCD) Type 1 (1 of 2)

SCD Type 1 (Overwrite Target with Latest Values)

- When a record updates, the **previous record is simply overwritten** by its **key** with the new value.
- When a record is deleted by its **key**, the **record is removed**
- There is **no tracking of old keys (rows)**, only the **current data** is retained.

customers (*target table with all customers*)

CustomerID	Name	Address	ProcessDate
1	Peter	1 Blue Rd.	5/1/2025
2	Samarth	22 Front St	5/1/2025



Update the **customers** table with the **new customer** information

updates (*source table with updates*)

CustomerID	Name	Address	Operation	ProcessDate
1	Peter	99 Sunny St..	UPDATE	5/15/2025
1	Peter	123 Main St.	UPDATE	5/20/2025
2	null	null	DELETE	5/20/2025
3	Kostas	100 Athens Dr.	NEW	5/20/2025



Change Data Capture (CDC) Overview

Slowing Changing Dimensions (SCD) Type 1 (2 of 2)

customers (*target table with all customers*)

Update the row

Delete the row

Insert the row

CustomerID	Name	Address	ProcessDate
1	Peter	1 Blue Rd.	5/1/2025
2	Samarth	22 Front St	5/1/2025
3	Kostas	100 Athens Dr.	5/20/2025



updates (*source table with updates*)

CustomerID	Name	Address	Operation	ProcessDate
1	Peter	99 Sunny St..	UPDATE	5/15/2025
1	Peter	123 Main St.	UPDATE	5/20/2025
2	null	null	DELETE	5/20/2025
3	Kostas	100 Athens Dr.	NEW	5/20/2025

Updated SCD Type 1 Table

CustomerID	Name	Address	ProcessDate
1	Peter	123 Main St.	5/20/2025
3	Kostas	100 Athens Dr.	5/20/2025

- **Peter** was updated with the **latest** information
- **Samarth** was deleted from the table
- **Kostas** was added to the table

Used **SCD Type 1** when only the **current data matters** and historical values are irrelevant



Change Data Capture (CDC) Overview

Slowing Changing Dimensions (SCD) Type 2 (Outside the scope of the

SCD Type 2 (Historical Tracking/Versioning)

- *When a record changes:*
 - The **old record is preserved** with an additional column indicating its validity period (start date, end date, or a current flag) and a **new record is inserted** with the updated information.
 - When a **record is marked as deleted**, the record is kept and a column indicates the record is **inactive**.
- Used when **historical data is important**, and the system needs to track **how attributes change over time** (like tracking changes in a customer's address or status).

Example customers target table with SCD Type 2

CustomerID	Name	Address	ProcessDate	__START_AT	__END_AT
1	Peter	123 Main St.	5/20/2025	5/20/2025	null
1	Peter	1 Blue Rd.	5/1/2025	5/1/2025	5/20/2025
2	Samarth	22 Front St	5/1/2025	5/1/2025	5/20/2025
3	Kostas	100 Athens Dr.	5/20/2025	5/20/2025	null

Columns indicate **active** and **inactive** rows

Null values indicate the **current row**

Non null values indicate **inactive (historic)** rows



Change Data Capture (CDC) Overview

Using AUTO CDC INTO in Spark Declarative Pipelines (Formerly APPLY CHANGES INTO)

customers (target table with all customers)

CustomerID	Name	Address	ProcessDate
1	Peter	1 Blue Rd.	5/1/2025
2	Samarth	22 Front St	5/1/2025



updates (source table with updates)

CustomerID	Name	Address	Operation	ProcessDate
1	Peter	99 Sunny St..	UPDATE	5/15/2025
1	Peter	123 Main St.	UPDATE	5/20/2025
2	Samarth	22 Front St	DELETE	5/20/2025
3	Kostas	100 Athens Dr.	NEW	5/20/2025

```
CREATE OR REFRESH STREAMING TABLE customers;
```

Create a streaming table, then use AUTO CDC to populate it

```
CREATE FLOW scd_type_1_flow AS
```

Apply updates, inserts and deletes to the **target** table

```
AUTO CDC INTO customers
```

Source records to determine **updates, deletes and inserts**

```
FROM STREAM updates
```

The column(s) that **uniquely identify a row** in the source and target tables

```
KEYS (CustomerID)
```

```
APPLY AS DELETE WHEN operation = "DELETE"
```

Specifies when a CDC event should be treated as a **DELETE** rather than an upsert.

```
SEQUENCE BY ProcessDate
```

Specifies the **logical order of CDC events** in the source data

```
COLUMNS * EXCEPT (operation)
```

Specifies a **subset of columns** to include in the **target**

```
STORED AS SCD TYPE 1;
```

Whether to store records as **SCD type 1 (default) or type 2**


*INSERTs and UPDATEs are implicitly implemented using the **KEY(s)**, no coding required*

Use **AUTO CDC INTO** (Formerly **APPLY CHANGE INTO**) instead of a traditional batch job with complex **MERGE INTO** syntax

Use **AUTO CDC FROM SNAPSHOT** (Public Preview as of 2025Q2) to process changes in database snapshots.



Change Data Capture (CDC) Overview






Building Lakeflow Spark Declarative Pipelines

DEMONSTRATION

Change Data Capture with AUTO CDC INTO INTO



Notebook: 6 Demo – Change Data Capture with AUTO CDC with SCD TYPE 1



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](https://www.apache.org/).



Building Lakeflow Spark Declarative Pipelines

LECTURE

Summary and Next Steps

Open **Notebook 7 – Summary and Next Steps** to explore additional topics and recommended next steps beyond this course.

Notebook: 7 – Summary and Next Steps





Building Lakeflow Spark Declarative Pipelines

LAB EXERCISE

BONUS – AUTO CDC INTO with SCD Type 1

If you are interested in performing CDC with Spark Declarative Pipelines we have provided an extra lab you can complete after class.



Notebook: 8 Lab – AUTO CDC INTO with SCD Type 1



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

